

Hadoop BDH Connector

Salesforce data at 150M+ scale via the BDH data mart — filter-first, with a 24-hour lag.

What it is

Reads a nightly mirror of Salesforce held on the BDH Hadoop data mart (cheaper than querying Salesforce live). Because it can see 150M+ rows, its defining feature is a mandatory **filter-first** design with a fail-closed guard, so it only ever scans the data you have deliberately selected.

Property	Detail
Source	BDH Hadoop data mart — nightly Salesforce mirror, ~24h behind
Connects with	WebHDFS (or a mounted local path); read-only HDFS principal
ACL model	owner / group / public — coarser than live Salesforce sharing (a documented trade-off)
Identity	Item id = Salesforce record Id → must use a separate Graph connection from the live Salesforce connector
Runtime	Standalone .NET 10 Windows service · 650 automated tests

Signature capabilities

- **Filter-first:** partition pruning (skip dt=/region= folders) + 14 record-predicate operators + per-object row cap
- **Fail-closed scale guard:** an object with no effective filter refuses to crawl (would scan 150M rows) unless explicitly allowed
- 24-hour dt-watermark drives incrementals after the nightly sync lands
- Delegation token never logged; WebHDFS retry ladder; owner/group/public ACL

Before setup — key values to request

Request these (see the full guide for the complete checklist)

Value / setting	Ask from
WebHDFS endpoint (incl. /webhdfs/v1) or the mount path	Hadoop / BDH platform team
A read-only HDFS principal (and private CA bundle if TLS-inspected)	Hadoop / BDH platform team
The Hive partition layout and object-to-path mapping	Hadoop / BDH platform team
A separate Graph connection id + sign-off on the per-object filter policy	Connector operator + data owner

Key signals to monitor

Metric / signal	Healthy looks like
Rows scanned vs matched vs pruned ratio	Most partitions pruned; scan is a small fraction of 150M
Guard-refusal count	Zero in steady state; a spike means an object lost its filter
Watermark age	Under ~26h; older means the nightly sync is late
..._dead_letter_depth	Near zero

Top troubleshooting

Symptom	First action
An object is refused ("full-scan refused")	Add a real filter in filters.json; use ALLOW_FULL_SCAN only if genuinely intended
Partial crawl / deletion sweep suppressed	An oversized file was skipped — handle that file; the sweep self-protects
Identity directory incomplete	Loud by design — the sync stops rather than apply coarse ACLs; fix the directory export
Watermark stale > 26h	Check the nightly BDH→Hadoop sync job upstream

Hadoop BDH Connector

Full detail: **05_Hadoop_BDH_Connector_Guide.docx** and the Tenant & Common Concepts Guide (00). Metric prefix shown is per-connector.